

# Weekly Assignment 6

Brandon Owens and Loan Pham

## Q. 1

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.svm import SVC
```

```
In [2]: # Import the dataset "iowa_housing.csv". The dataset is about housing prices in Ames, Iowa.
# You are going to predict sale price.
df = pd.read_csv("../dataFiles/iowa_housing.csv")
df.head()
```

```
Out[2]:
```

	<b>Id</b>	<b>MSSubClass</b>	<b>MSZoning</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>Street</b>	<b>Alley</b>	<b>LotShape</b>	<b>LandContour</b>	<b>Utilities</b>	<b>SalePrice</b>
<b>0</b>	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	1633900
<b>1</b>	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	716000
<b>2</b>	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	1500000
<b>3</b>	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	1815000
<b>4</b>	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	2333000

5 rows × 81 columns



```
In [3]: # (a) Show the head and shape of the dataframe. Drop "Id" in the dataset.
df_drop = df.drop(["Id"], axis=1)
df_drop.head()
```

```
Out[3]:
```

	<b>MSSubClass</b>	<b>MSZoning</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>Street</b>	<b>Alley</b>	<b>LotShape</b>	<b>LandContour</b>	<b>Utilities</b>	<b>SalePrice</b>
<b>0</b>	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	1633900
<b>1</b>	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	716000
<b>2</b>	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	1500000
<b>3</b>	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	1815000

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotC
4	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	

5 rows × 80 columns



In [4]: `df_drop.shape`


Out[4]: (1460, 80)

In [5]: `# (b) Use the dataset to:  
# (i) create a correlation matrix (df.corr()), and then draw the heatmap on the  
corr = df_drop.corr()  
corr.head()`

Out[5]:


	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
<b>MSSubClass</b>	1.000000	-0.386347	-0.139781	0.032628	-0.059316	0.027850	0.040581
<b>LotFrontage</b>	-0.386347	1.000000	0.426095	0.251646	-0.059213	0.123349	0.088866
<b>LotArea</b>	-0.139781	0.426095	1.000000	0.105806	-0.005636	0.014228	0.013788
<b>OverallQual</b>	0.032628	0.251646	0.105806	1.000000	-0.091932	0.572323	0.550684
<b>OverallCond</b>	-0.059316	-0.059213	-0.005636	-0.091932	1.000000	-0.375983	0.073741

5 rows × 37 columns




In [6]: `sns.heatmap(corr)`

Out[6]: <AxesSubplot:>



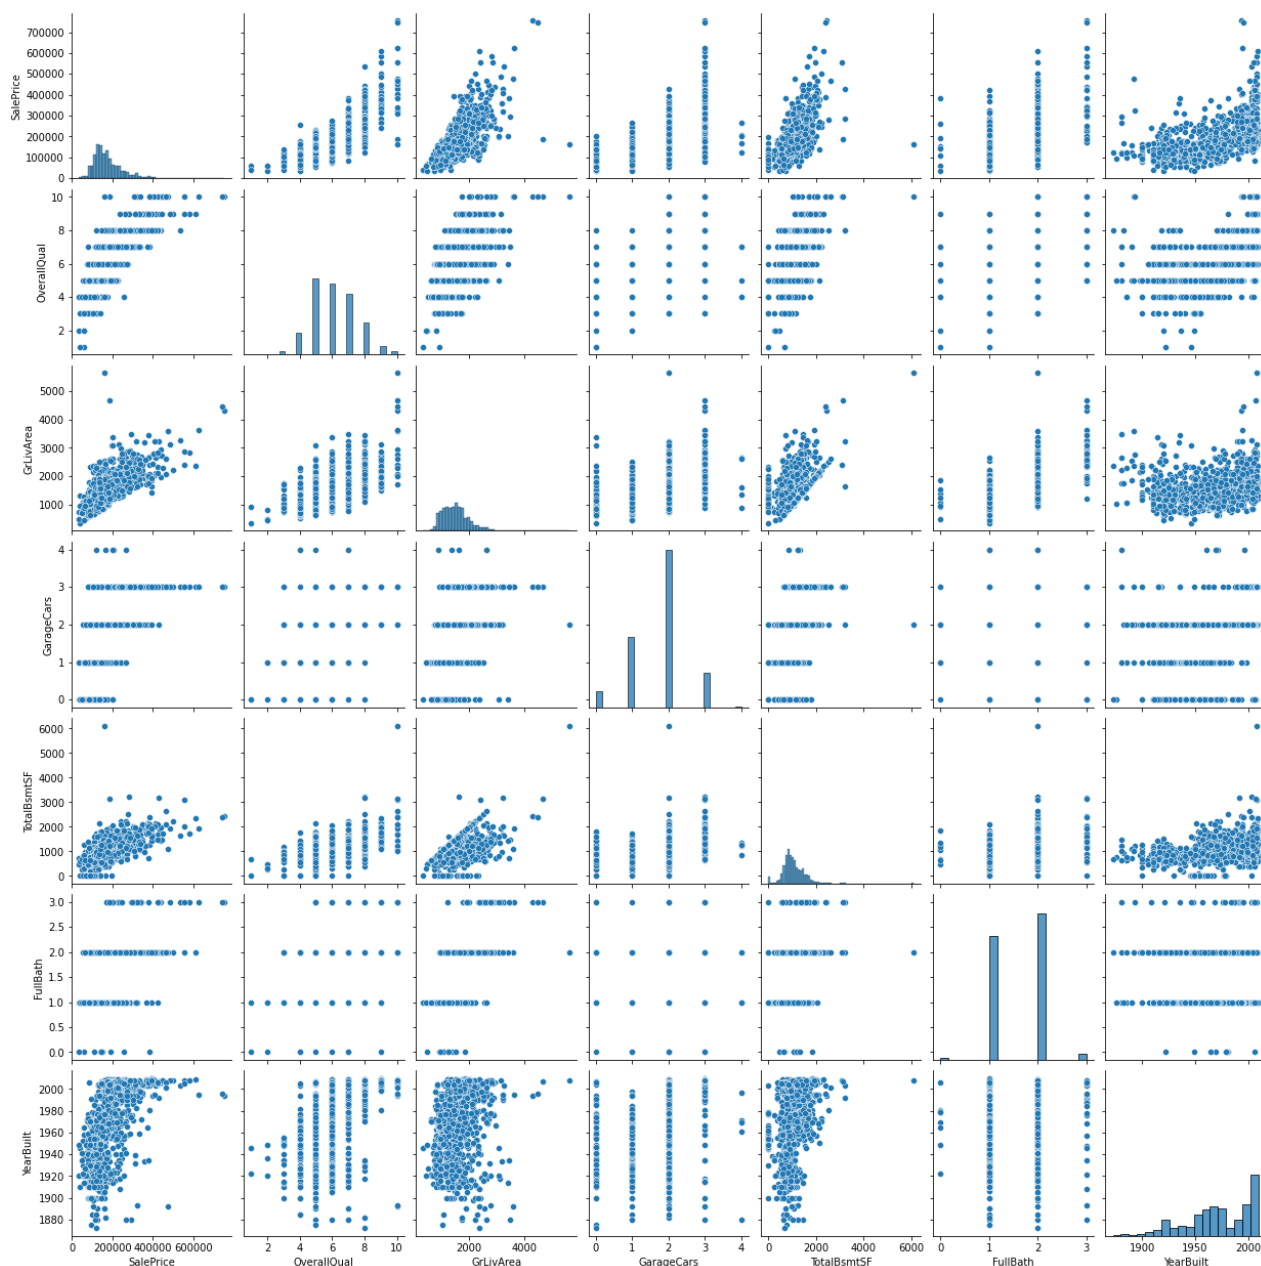
```
In [7]: # (ii) Find the features that the absolute value of correlation with "SalePrice"
# Then, draw the heatmap of correlation on those features.
features_abs = corr[abs((corr.SalePrice)>=.5)].SalePrice.keys()
abs_corr = corr.loc[features_abs, features_abs]
sns.heatmap(abs_corr)
```

Out[7]: <AxesSubplot:>



```
In [8]: # (iii) create pair plots on the following features
sns.pairplot(df_drop, vars=['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars',
```

Out[8]: <seaborn.axisgrid.PairGrid at 0x22200026ac0>



```
In [9]: # (iv) Find the most important feature relative to the "SalePrice" based on absolute correlation
corr.sort_values(["SalePrice"], ascending = False, inplace = True)
corr.SalePrice.head()
```


```
Out[9]: SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
Name: SalePrice, dtype: float64
```

```
In [10]: # (v) create a distribution plot on "SalePrice"
sns.distplot(df_drop['SalePrice'], kde = True)
```

C:\Users\jeric\miniconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: 'distplot' is a deprecated function and will be removed in a future version. Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[10]: <AxesSubplot:xlabel='SalePrice', ylabel='Density'>




```
In [11]: # (vi) For all the numerical features except "SalePrice",
# replace all the missing values of numerical features with the median value of ea
temp=pd.DataFrame()
temp['SalePrice'] = df['SalePrice']
num_cols = df.select_dtypes(exclude = ['object']).columns
num_cols = num_cols.drop('SalePrice')
med_replace = df[num_cols]
med_replace = med_replace.fillna(med_replace.median())
med_replace
```

Out[11]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	N
0	1	60	65.0	8450	7	5	2003	2003	
1	2	20	80.0	9600	6	8	1976	1976	
2	3	60	68.0	11250	7	5	2001	2002	
3	4	70	60.0	9550	7	5	1915	1970	
4	5	60	84.0	14260	8	5	2000	2000	
...	...	...	...	...	...	...	...	...	
1455	1456	60	62.0	7917	6	5	1999	2000	
1456	1457	20	85.0	13175	6	6	1978	1988	
1457	1458	70	66.0	9042	7	9	1941	2006	
1458	1459	20	68.0	9717	5	6	1950	1996	
1459	1460	20	75.0	9937	5	6	1965	1965	

1460 rows × 37 columns




```
In [12]: # (vii) Create dummies for all categorical features.
# The final shape of dataset should be (1460, 246) (set drop_first = True)
dummies_cols = df.select_dtypes(include = ['object']).columns
```

```
new_housing = df[dummies_cols]
new_housing = pd.get_dummies(new_housing, drop_first = True)
new_housing.head()
```

```
Out[12]:
```

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave	Alley_Pave	LotShape_IR2	
0	0	0	1	0	1	0	0	
1	0	0	1	0	1	0	0	
2	0	0	1	0	1	0	0	
3	0	0	1	0	1	0	0	
4	0	0	1	0	1	0	0	

5 rows × 209 columns



```
In [13]: newdf_housing = pd.concat([temp, med_replace, new_housing], axis =1)
newdf_housing.shape
```

Out[13]: (1460, 247)

```
In [14]: newdf_housing['SalePrice']
```

```
Out[14]:
```

0	208500
1	181500
2	223500
3	140000
4	250000
...	
1455	175000
1456	210000
1457	266500
1458	142125
1459	147500

Name: SalePrice, Length: 1460, dtype: int64

```
In [15]: # (c) Then, do the modelling:
```

```
In [16]: # (i) Check for any missing values again.
newdf_housing.isnull().any()
```

```
Out[16]:
```

SalePrice	False
Id	False
MSSubClass	False
LotFrontage	False
LotArea	False
...	
SaleCondition_AdjLand	False
SaleCondition_Alloca	False
SaleCondition_Family	False
SaleCondition_Normal	False
SaleCondition_Partial	False


Length: 247, dtype: bool

```
In [17]: # (ii) create y = "SalePrice". Drop y from the dataframe.
y = "SalePrice"
house_temp = newdf_housing.drop([y], axis =1)
house_temp.head()
```

```
Out[17]:
```

	<b>Id</b>	<b>MSSubClass</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>OverallQual</b>	<b>OverallCond</b>	<b>YearBuilt</b>	<b>YearRemodAdd</b>	<b>MasVnr</b>
<b>0</b>	1	60	65.0	8450	7	5	2003	2003	
<b>1</b>	2	20	80.0	9600	6	8	1976	1976	
<b>2</b>	3	60	68.0	11250	7	5	2001	2002	
<b>3</b>	4	70	60.0	9550	7	5	1915	1970	
<b>4</b>	5	60	84.0	14260	8	5	2000	2000	

5 rows × 246 columns



```
In [18]: # (iii) Split the dataset into train and test.

X_train, X_test, y_train, y_test = train_test_split(house_temp, newdf_housing, test
print("X_train : " + str(X_train.shape))
print("X_test : " + str(X_test.shape))
print("y_train : " + str(y_train.shape))
print("y_test : " + str(y_test.shape))
```


```
X_train : (1022, 246)
X_test : (438, 246)
y_train : (1022, 247)
y_test : (438, 247)
```

```
In [19]: # (iv) Train a linear regression model.
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[19]: LinearRegression()
```

```
In [20]: # (v) predict "SalePrice" with the test data.
y_pred = model.predict(X_test)
plt.scatter(y_test, y_pred)
plt.xlabel('Y Test (True Values)')
plt.ylabel('Predicted Values')
plt.title('Predicted vs. Actual Values ')
```

```
Out[20]: Text(0.5, 1.0, 'Predicted vs. Actual Values ')
```



```
In [21]: # (vi) find the root mean squared error of the model on the test data.
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("Root Mean Squared Error = ", rmse)
print("The range of temperature = ", y_test.min(), y_test.max())
```

```
Root Mean Squared Error = 2681.0955819197743
The range of temperature = SalePrice          35311.0
Id          11.0
MSSubClass  20.0
LotFrontage 21.0
LotArea     1491.0
...
SaleCondition_AdjLand 0.0
SaleCondition_Alloca  0.0
SaleCondition_Family  0.0
SaleCondition_Normal  0.0
SaleCondition_Partial 0.0
Length: 247, dtype: float64 SalePrice          755000.0
Id          1455.0
MSSubClass  190.0
LotFrontage 174.0
LotArea     70761.0
...
SaleCondition_AdjLand 0.0
SaleCondition_Alloca  1.0
SaleCondition_Family  1.0
SaleCondition_Normal  1.0
SaleCondition_Partial 1.0
Length: 247, dtype: float64
```

## Q2

```
In [22]: # (a) Read in the dataset 'online_shoppers_intention.csv'
df_shoppers = pd.read_csv('../dataFiles/online_shoppers_intention.csv')
df_shoppers.head()
```


```
Out[22]:   Administrative  Administrative_Duration  Informational  Informational_Duration  ProductRelated  Product_
0          0.0          0.0          0.0          0.0          1.0
```



	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration
1	0.0	0.0	0.0	0.0	2.0	0.0
2	0.0	-1.0	0.0	-1.0	1.0	0.0
3	0.0	0.0	0.0	0.0	2.0	0.0
4	0.0	0.0	0.0	0.0	10.0	0.0


```
In [23]: # (b) Count how many shoppers buy, i.e. "Revenue"==True.
# Count how many shoppers browse in the weekends, i.e. "Weekends"==True. Create the
fig, ax =plt.subplots(1,2)
df_shoppers['Revenue'].value_counts().plot.bar(ax=ax[0])
df_shoppers['Weekend'].value_counts().plot.bar(ax=ax[1])
```

Out[23]: <AxesSubplot:>




```
In [24]: # (c) Create the following plot, which shows the proportions of various kinds of
# (first, use value counts, and then use .plot.pie(autopct = '%.2f%%')
df_shoppers['VisitorType'].value_counts().plot.pie(autopct = '%.2f%%')
plt.legend()
```

Out[24]: <matplotlib.legend.Legend at 0x22202a77370>



```
In [25]: # (d) Check the month with most shoppers visiting the online shopping sites. Create a bar chart of the number of shoppers by month.
months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
new_df = df_shoppers['Month'].value_counts()
new_df = new_df.reindex(months)
new_df.plot.bar()
```

Out[25]: <AxesSubplot:>




```
In [26]: # (e) For all the numerical variables, fill the missing values with the median of the variable.
num_features = df_shoppers.select_dtypes(exclude = ["object"]).columns
num_shoppers = df_shoppers[num_features]
num_shoppers = num_shoppers.fillna(num_shoppers.median())
num_shoppers
```

Out[26]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated
0	0.0	0.0	0.0	0.0	1.0
1	0.0	0.0	0.0	0.0	2.0
2	0.0	-1.0	0.0	-1.0	1.0
3	0.0	0.0	0.0	0.0	2.0
4	0.0	0.0	0.0	0.0	10.0
...	...	...	...	...	...
12325	3.0	145.0	0.0	0.0	53.0
12326	0.0	0.0	0.0	0.0	5.0
12327	0.0	0.0	0.0	0.0	6.0
12328	4.0	75.0	0.0	0.0	15.0
12329	0.0	0.0	0.0	0.0	3.0

12330 rows × 6 columns



```
In [27]: # (f) Create dummies for categorical variables.
```

```
cate_vari = df_shoppers.select_dtypes(include = ["object"]).columns
shop_cat = df_shoppers[cate_vari]
shop_cat = pd.get_dummies(shop_cat, drop_first = True)
shop_cat.head()
```

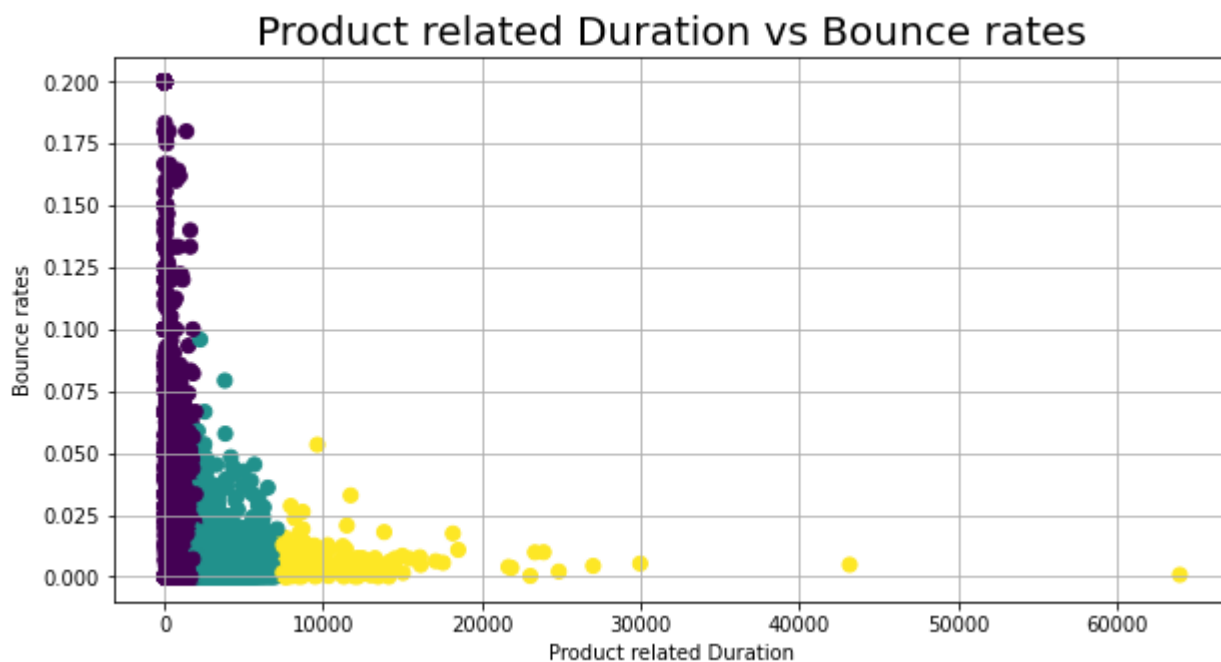
```
Out[27]:
```

	Month_Dec	Month_Feb	Month_Jul	Month_June	Month_Mar	Month_May	Month_Nov	Month_Oct
0	0	1	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0
2	0	1	0	0	0	0	0	0
3	0	1	0	0	0	0	0	0
4	0	1	0	0	0	0	0	0

```
In [28]: # (g) Use KMeans to group customers into 3 clusters. That is unsupervised learning
```

```
kmeans = KMeans(n_clusters=3)
kmeans.fit(num_shoppers)
y_kmeans = kmeans.predict(num_shoppers)
plt.figure(figsize=(10,5))
plt.title('Product related Duration vs Bounce rates', fontsize = 20)
plt.grid()
plt.xlabel('Product related Duration')
plt.ylabel('Bounce rates')
plt.scatter(df_shoppers["ProductRelatedDuration"], df_shoppers["BounceRates"], c=
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x222056b79d0>
```



```
In [29]: # (h) Set y = "Revenue" and X is the dataframe without "Revenue".
numerical_cols = df_shoppers.select_dtypes(exclude = ["number"]).columns
shop_nums = df_shoppers[numerical_cols]
df_dummies = pd.get_dummies(shop_nums, drop_first=True)
y = df_dummies["Revenue"]
```

```
X = df_dummies.drop("Revenue", axis=1)
df_dummies.isnull().any()
```

```
Out[29]: Weekend      False
Revenue      False
Month_Dec    False
Month_Feb    False
Month_Jul    False
Month_June   False
Month_Mar    False
Month_May    False
Month_Nov    False
Month_Oct    False
Month_Sep    False
VisitorType_Other      False
VisitorType_Returning_Visitor  False
dtype: bool
```

```
In [30]: # (i) Split the dataset into train and test sets.
# Train a support vector machine classifier model. Use kernel = "rbf" and class_weight = "balanced"
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
model = SVC(kernel='rbf', class_weight='balanced')
model.fit(X_train, y_train)
```

```
Out[30]: SVC(class_weight='balanced')
```

```
In [31]: # (j) Predict which online shopper will do a purchase. Find the accuracy score.
y_pred = model.predict(X_test)
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

```
Out[31]: 0.6387318751536004
```

### Q3

```
In [32]: # (a) Load the dataset.
from sklearn.datasets import fetch_olivetti_faces
face = fetch_olivetti_faces()
faces = fetch_olivetti_faces().images
```

```
In [33]: # (b) Show the first 10 images
fig, ax = plt.subplots(1, 10, figsize=(64, 64))
for i, axi in enumerate(ax.flat):
    axi.imshow(faces[i], cmap=plt.cm.bone)
    axi.set(xticks=[], yticks=[])
```



```
In [34]: # (c) Size of each image is 64x64. Use PCA to reduce it into 90 features.
# For the first row, show the first 10 original images.
# Then for the second row, show the first 10 images with reduced number of features.
x = face.data
```


```

y = face.target
pca = PCA(n_components = 90)
pca.fit(x)
transformed_data = pca.fit_transform(x)
x_approx = pca.inverse_transform(transformed_data)
x_approx_img = x_approx.reshape(400,64,64)

fig, ax = plt.subplots(1, 10, figsize=(64, 64))
for i, axi in enumerate(ax.flat):
    axi.imshow(faces[i], cmap=plt.cm.bone)
    axi.set(xticks=[], yticks=[])

fig, ax = plt.subplots(1, 10, figsize=(64, 64))
for i, axi in enumerate(ax.flat):
    axi.imshow(x_approx_img[i], cmap = plt.cm.bone)
    axi.set(xticks=[], yticks=[])

```



```

In [35]: # (d) Using images of reduced features to conduct the machine learning task.
# (i) Split the dataset into train and test.

X_train, X_test, y_train, y_test = train_test_split(face.data, face.target, random_

```

```

In [36]: pca = PCA(n_components=90)
pca.fit(X_train)

```

```

Out[36]: PCA(n_components=90)

```

```

In [37]: # (ii) Train a Random Forest Classifier model. Set n_estimators=100. Predict the

model = RandomForestClassifier(n_estimators=100)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

```

```

In [38]: # (iii) Find the accuracy score.
accuracy_score(y_test, y_pred)

```

```

Out[38]: 0.95

```

```

In [39]: # (iv) Create a confusion matrix, and put it in a heatmap. (Your heatmap may look
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

```

```

In [40]: sns.heatmap(cm)

```

Out[40]: &lt;AxesSubplot:&gt;

